







(a) GPT-4-Turbo




(b) Claude-Sonnet



(c) DeepSeek-R1



(d) Gemini-Flash



(e) Qwen-72B

Interpretation:

Lower points indicate better stability compared to the random baseline.

Random Baseline (Bar)
 AI Agent Performance (Point)